

Chitresh Yadav

Backend & Applied ML Engineer

+91 91097 31314 | chitreshy20@gmail.com | [linkedin.com/in/csy20](https://www.linkedin.com/in/csy20) | github.com/csy20 | csy20.me

SUMMARY

Software engineer focused on backend services and applied ML systems. Built a streaming speech-to-speech pipeline (ASR → MT → TTS over WebRTC) with measured 2505 ms mean / 2580 ms p50 end-to-end latency, then cut latency 23.6% with Silero-VAD segmentation. Fine-tuned Whisper with LoRA for Hindi-English code-mixed ASR and designed an ONNX export path for Rust inference. Also shipped queue-based media workers (Redis, PostgreSQL, MinIO, Docker Compose) and a token-efficient LLM router with OpenTelemetry. Seeking SDE-1 / Backend / Applied ML roles.

SKILLS

Languages: Python, C++, Rust, TypeScript / Node.js, SQL

Backend: FastAPI, Express.js, REST APIs, WebRTC (aiortc), asyncio, Redis queues, PostgreSQL, SQLite, MongoDB, MinIO / S3

ML / AI: Whisper / faster-whisper, LoRA / PEFT, ONNX Runtime (fp16 / int8), IndicTrans2, Silero-VAD, BART, LLM routing, MCP, WER / CER eval, OpenTelemetry

Infra: Docker, Docker Compose, GitHub Actions, GCP (Cloud Run, Cloud Storage, Firebase), Linux

PROJECTS

speech-relay — Real-Time Speech-to-Speech Translation github.com/csy20/speech-relay

Python · faster-whisper · IndicTrans2 · Coqui / edge-tts · aiortc · ONNX Runtime · asyncio

- Built a streaming ASR → MT → TTS service that returns live translated audio over WebRTC (EN→HI), instrumented per hop at 2504.73 ms mean / 2579.7 ms p50 end-to-end.
- Replaced fixed-window flushing with Silero-VAD segmentation and tuned buffers, cutting end-to-end latency by 23.61%.
- Designed a fan-out path so one shared ASR stream serves N listeners with independent translation targets (1.56× parallel efficiency at 2 listeners).

speech-relay-rust — Code-Mixed ASR Fine-Tune + ONNX Serving github.com/csy20/speech-relay-rust

Python · Whisper · LoRA / PEFT · PyTorch · ONNX · Rust (ort) · MUCS / OpenSLR-104

- Implemented a data pipeline over MUCS Hindi-English speech: 16 kHz resample, code-mixed script filter (Devanagari + Latin), train/val/test split, Hugging Face export.
- Fine-tuned Whisper with LoRA on decoder attention (q/k/v/out proj) so the model handles mid-utterance Hindi-English switches; training path runs on a T4 in a few hours.
- Designed the serving path: merge LoRA weights, export encoder/decoder to ONNX, int8 quantize, run inference from a Rust ONNX Runtime (ort) service outside Python.

MediaPipe AI — Distributed Media Processing Pipeline github.com/csy20/mediapipe-ai

Express · Python workers · Redis · PostgreSQL · MinIO · Docker Compose · Whisper · BART

- Architected a 5-container system: React client → Express REST gateway → Redis job queue → Python workers (Whisper transcription + BART summarization) → PostgreSQL metadata + MinIO object storage.
- Added async job status polling, structured logging, and S3-compatible uploads for audio/video up to 500 MB so workers scale independently of the API.

router-agent — Token-Efficient LLM Router github.com/csy20/router-agent

Python · Docker · Ollama · Fireworks API · OpenTelemetry / SigNoz · AMD Hackathon ACT II

- Built a hybrid router for the AMD Developer Hackathon that classifies tasks (math, NER, summarization, code, QA) with zero-token heuristics, then escalates to a local 1.5B LLM or at most one paid API call.

OPEN SOURCE

ripwire (Red Hat Emerging Technologies) — merged PR #244 github.com/redhat-et/ripwire/pull/244

C++23 · TypeScript / JS call-graph · MCP · static analysis · ~2.2k stars

- Merged a call-binding fix in Red Hat ET's ripwire (C++23 CLI + MCP server for coding agents): TS/JS literal receivers no longer pin unrelated same-named functions.
- On webpack, removed 139 false bindings (repeat/join/exec/filter) and lifted External matches 81 → 611; node lib/ gained 10 correct externals. Fixes #163. [PR #244](https://github.com/redhat-et/ripwire/pull/244).

EDUCATION

B.Tech, Computer Science and Business Systems

Shri Shankaracharya Technical Campus, Bhilai, Chhattisgarh

May 2026